Practical Introduction to HPC
Cellular Automaton Exercise

Aim

The aim of this exercise is to design serial and parallel algorithms to run a simple cellular automaton which attempts to model traffic flow. You should write some code to implement the serial algorithm

The Model

The simulation box is a line of N cells numbered $1, 2, \ldots, N-1, N$ (the road) which can each only have two values: 1 if a car is present on that section of road; 0 if it is empty. We will denote the value of cell number i at time t by $R^t(i)$.

The *update rules* for each time step are very simple:

- if the space in front of a car is empty then it moves forward one cell;

- otherwise it stays where it is.

Remember that we proceed through a series of iterations where the values of the cells at time $t+1$ depend entirely on the values at time t. In other words you should be able to apply the update rules to the cells at time t in any order and you should still get the same values at $t+1$.

After each iteration we compute the *average velocity* of the cars. This is calculated as the number of cars that move in an iteration divided by the total number of cars, and is a value between 0 and 1.

We use *periodic boundary conditions*, i.e. when a car moves off the right-hand-side of the road it reappears on the left, and vice-versa (this is the same as saying we are simulating a roundabout rather than a straight section of road). To do this, we identify $R^t(0)$ (the cell to the left of cell number 1) with $R^t(N)$ and $R^t(N + 1)$ (the cell to the right of cell number N) with $R^t(1)$.

1 Rules for Cellular Automaton

It should be fairly clear that, when we update to a new time step, the new value $R^{t+1}(i)$ depends on its old value at time t, and also on the old values of the neighbours. In other words, $R^{t+1}(i)$ depends on the three values $R^t(i - 1)$, $R^t(i)$ and $R^t(i + 1)$.

You should fill in Table 1 and Table 2 to get the explicit form of the update rules for this cellular automaton. For convenience, we have split the 8 possible cases into two tables, one for an initially empty cell $R^t(i) = 0$, and the other for a full cell $R^t(i) = 1$. Writing things out in this explicit form you can see that there are in fact 256 possible 1D cellular automata (each corresponding to different update rules), but very few of them are in fact of any interest!

	$R^t(i-1) = 0$	$R^t(i-1) = 1$
$R^t(i+1) = 0$		
$R^t(i+1) = 1$		

Table 1: Values of $R^{t+1}(i)$ if $R^t(i) = 0$

	$R^t(i-1) = 0$	$R^t(i-1) = 1$
$R^t(i+1) = 0$		
$R^t(i+1) = 1$		

Table 2: Value of $R^{t+1}(i)$ if $R^t(i) = 1$

2   Serial Program

Having written the update rules, **write some code (in any programming language)** that implements them for a road of length N. The easiest approach is to have two arrays, one corresponding to the *old* values at time t and the other to the *new* values at $t+1$. Loop through the cells computing the values of the the *new* array based on *old*. After each complete time step, copy the *new* array back to *old* in preparation for the next time step. You must consider how to implement the periodic boundary conditions.

We need to define some initial conditions, i.e. the state of the road at time $t = 0$. This is best done using a random number generator; such generators typically return a real value distributed evenly between 0.0 and 1.0. We need to generate an initial road with a specific density of cars, e.g. we might want a reasonably congested road with a density of 0.75 (where, on average, three out of four cells are occupied):

- loop over the whole road $i = 1, 2, \ldots, N - 1, N$
- pick a random number $rand$ in the range 0.0 to 1.0.
- if $rand < density$
    - set $R^0(i) = 1$
- else
    - set $R^0(i) = 0$
- end loop over i

It is easy to see that this works for densities of 0.0 and 1.0 (random number generators normally have the range $0.0 \leq rand < 1.0$), and in fact works for any value of the density.

2.1   Experiments

The aim is to reproduce the graph of average velocity versus car density that was outlined in the lectures.

- pick a large value for N, a few thousand should suffice
- choose a target value for the car density
- initialise the road as above
- compute the actual density which will be close to (but not exactly the same as) your target
- loop over many time steps $t = 0, 1, 2, \ldots$
    - update the road from time t to $t + 1$ by updating all the cells
    - compute the average car velocity for this time step
- end loop over time steps

You should do this for a range of densities to map out how the average velocity depends on density, and plot a graph of the results.

You can just use a single (large) value of N, although you could check that the results do not depend strongly on this value (e.g. run two simulations of size N and $2N$ for the same density).

Here are a few hints to help you:

1. As the initial distribution of cars is entirely random, the simulation may take many time steps to settle down to a steady state. You should see the velocity varying with time step, but eventually settling down to a fairly steady value. What you are looking for is the *asymptotic* velocity, i.e. the average velocity per step after many time steps.

2. The amount of time taken to reach this steady state depends on the density (and the road length N), so it may take some experimentation to find out how long to run the model.

3. To check your code is correct, try a few simple examples where you know the answer:

   - a solid road (density of 1.0) should have velocity 0.0

   - a solitary car should move forward with velocity 1.0

   - if you fill half of the road with cars in a line and leave the other half free, the average velocity should increase linearly with time step until it reaches a maximum of 1.0 after $N/2$ time steps (think of the cars at a set of traffic lights)

   - the number of cars should always remain the same